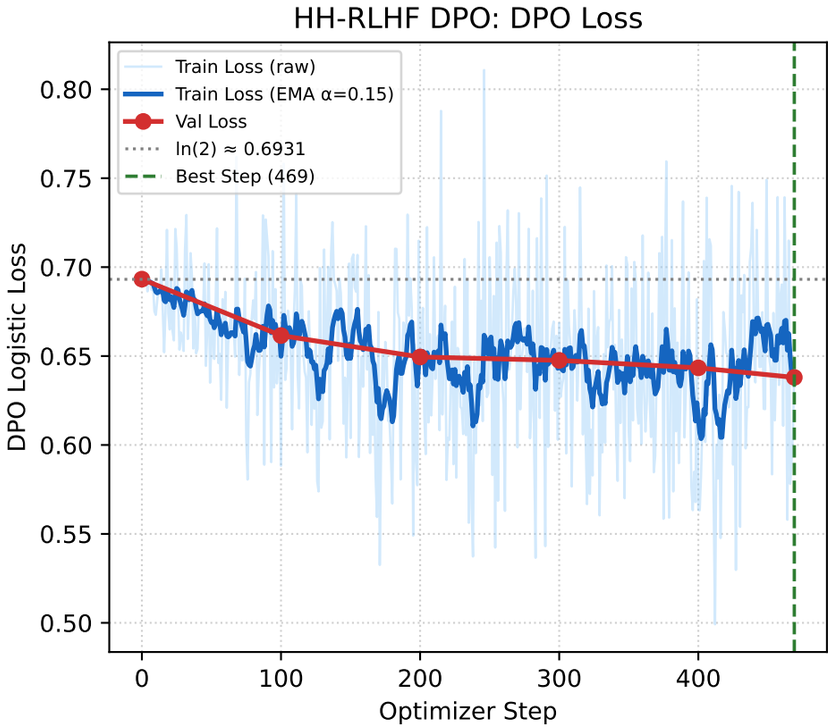
HH-RLHF DPO: DPO Loss
Train Loss (raw)
Train Loss (EMA α=0.15)
Val Loss
ln(2) ≈ 0.6931
Best Step (469)
DPO Logistic Loss
Optimizer Step

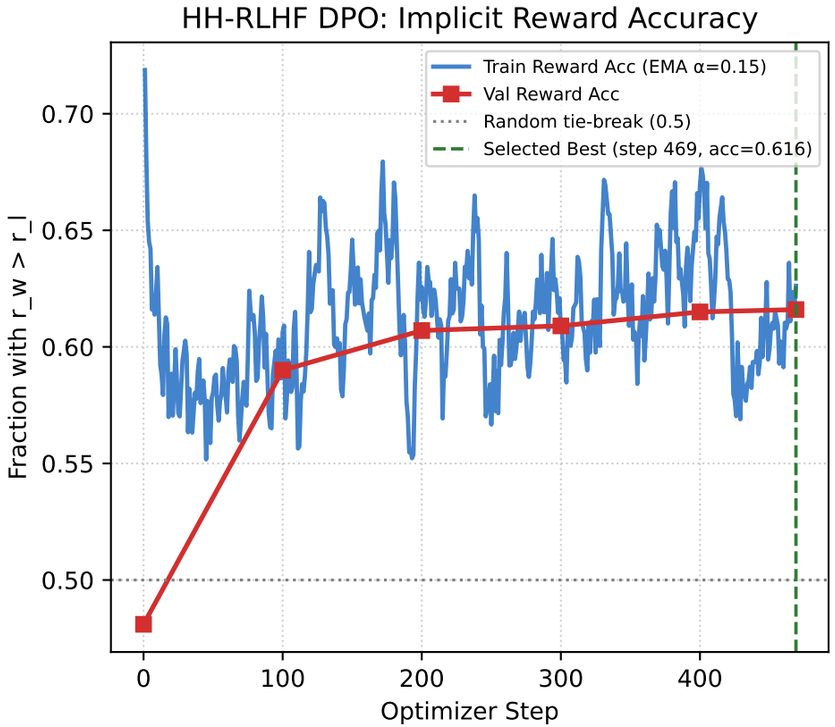
HH-RLHF DPO: Implicit Reward Accuracy
Train Reward Acc (EMA α=0.15)
Val Reward Acc
Random tie-break (0.5)
Selected Best (step 469, acc=0.616)
Fraction with r_w > r_l
Optimizer Step

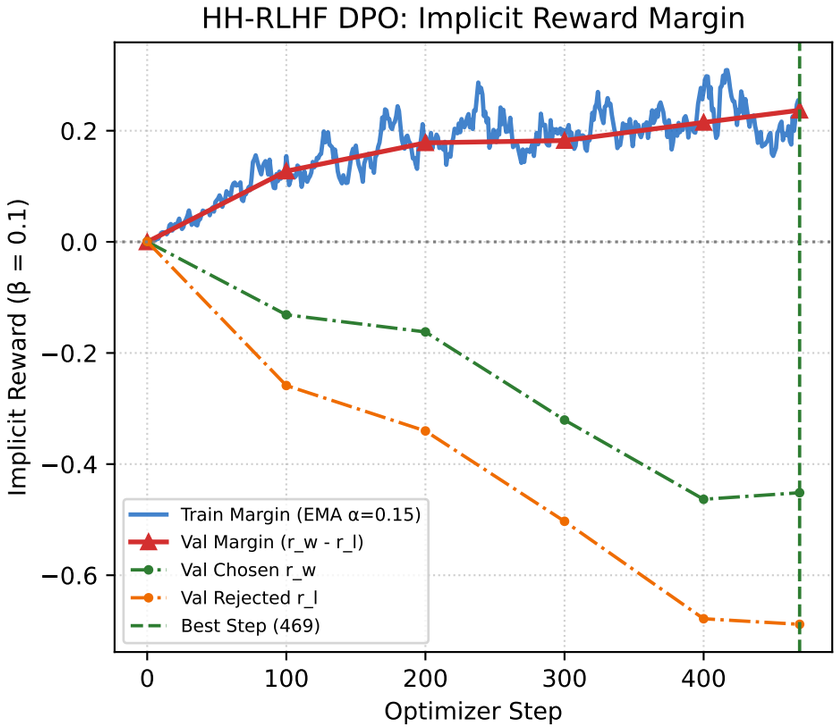
HH-RLHF DPO: Implicit Reward Margin
Train Margin (EMA α=0.15)
Val Margin (r_w - r_l)
Val Chosen r_w
Val Rejected r_l
Best Step (469)
Implicit Reward (β = 0.1)
Optimizer Step