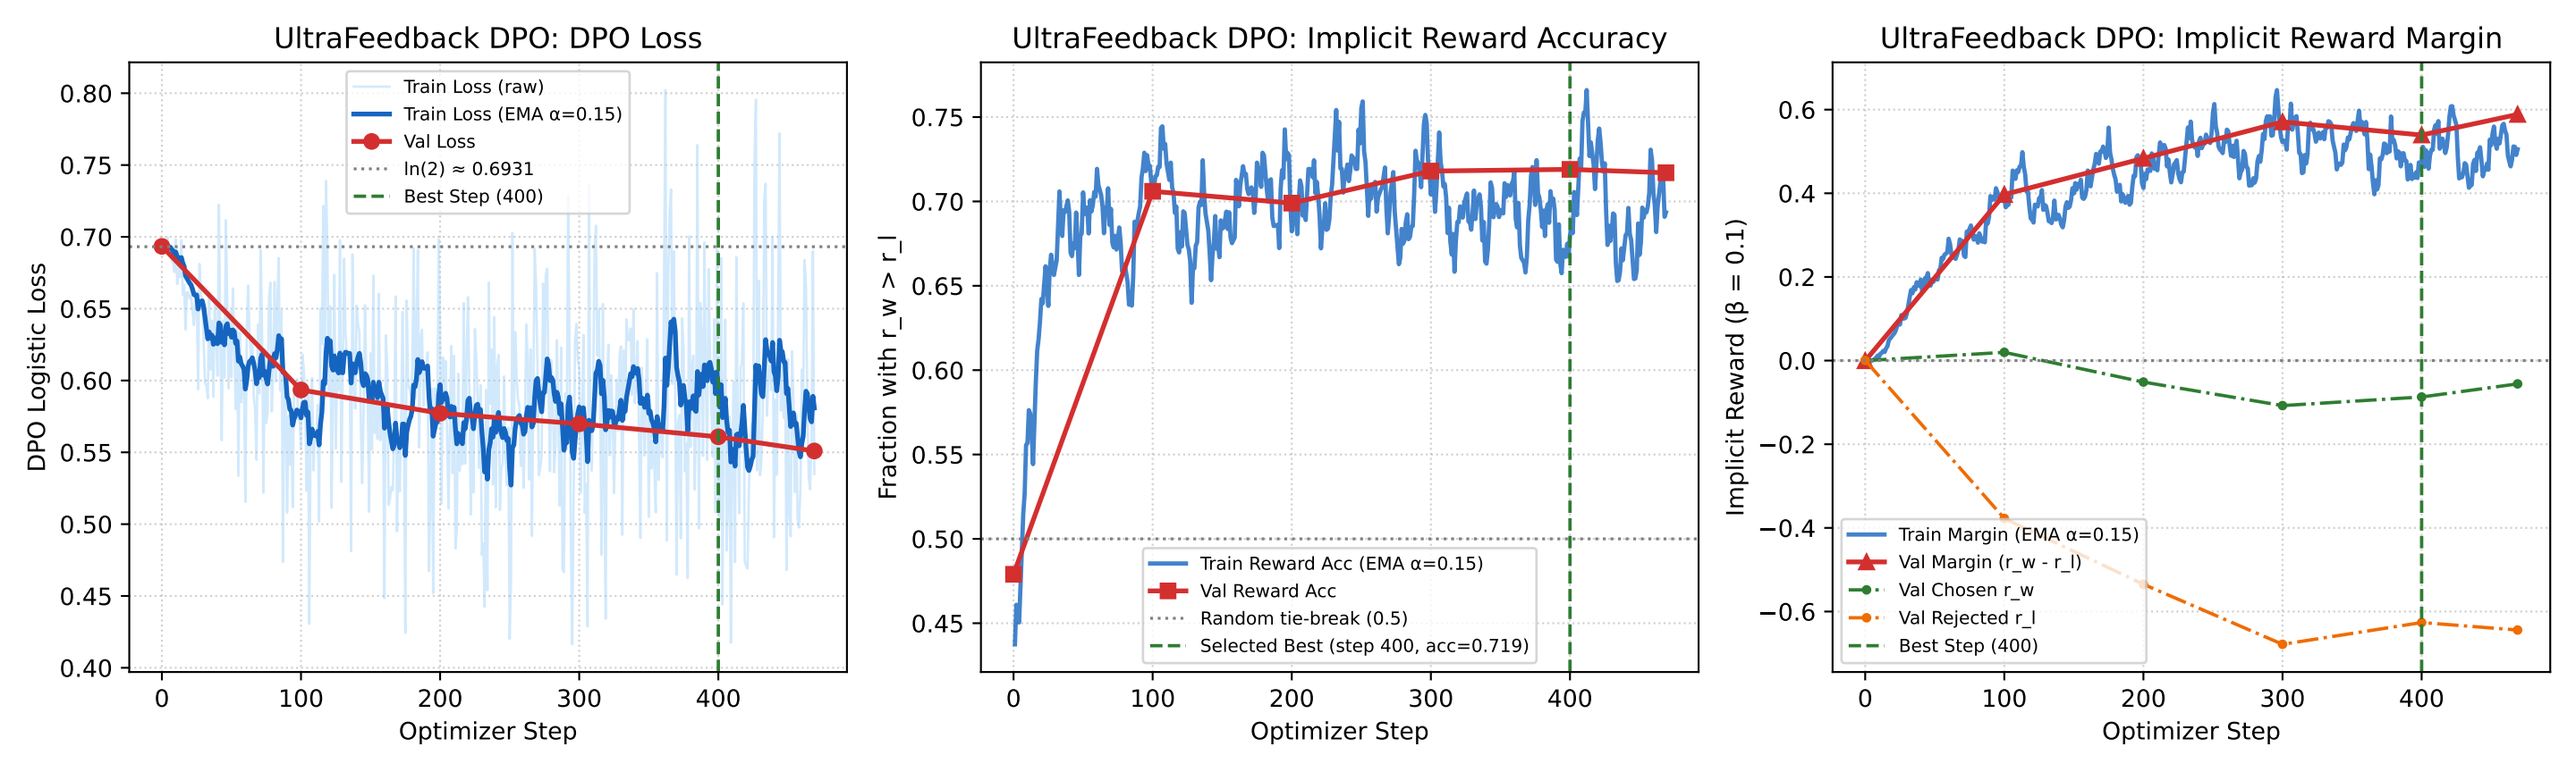
UltraFeedback DPO: DPO Loss
DPO Logistic Loss
Optimizer Step
Train Loss (raw)
Train Loss (EMA α=0.15)
Val Loss
ln(2) ≈ 0.6931
Best Step (400)
UltraFeedback DPO: Implicit Reward Accuracy
Fraction with r_w > r_l
Optimizer Step
Train Reward Acc (EMA α=0.15)
Val Reward Acc
Random tie-break (0.5)
Selected Best (step 400, acc=0.719)
UltraFeedback DPO: Implicit Reward Margin
Implicit Reward (β = 0.1)
Optimizer Step
Train Margin (EMA α=0.15)
Val Margin (r_w - r_l)
Val Chosen r_w
Val Rejected r_l
Best Step (400)