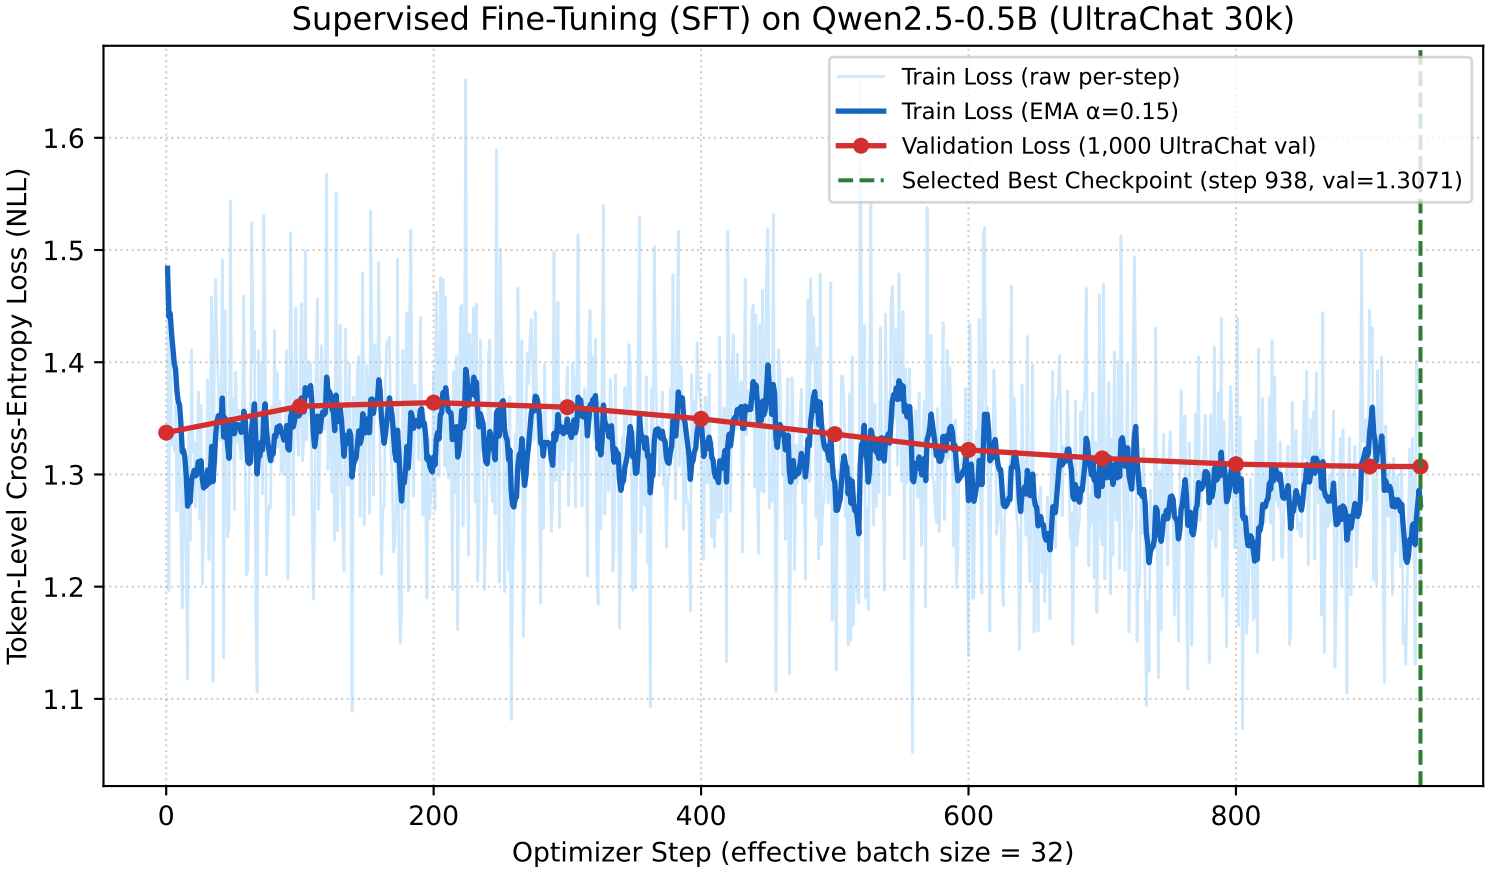
Supervised Fine-Tuning (SFT) on Qwen2.5-0.5B (UltraChat 30k)
Train Loss (raw per-step)
Train Loss (EMA α=0.15)
Validation Loss (1,000 UltraChat val)
Selected Best Checkpoint (step 938, val=1.3071)
Token-Level Cross-Entropy Loss (NLL)
1.6
1.5
1.4
1.3
1.2
1.1
0
200
400
600
800
Optimizer Step (effective batch size = 32)